



Case Study: Kubernetes Deployment & Service

Developed By: Ashoka Naidu, Devanshu Sharma, Hemanth Singh Kumar, Molshree Singhal, & Shushovan Chakraborty

Table of Contents

[Table of Contents](#)

[Deployment](#)

[Deployment Controller changes the actual state to the desired state at a controlled rate. We can create new ReplicaSets, remove existing Deployments and adopt all their resources with new Deployments.](#)

[Step 1: Create Deployment File](#)

[Create a deployment.yml in a vi editor on your master machine](#)

[Get into the deployment.yml file](#)

[Step 2: Deploy it & verify it](#)

[Service](#)

[Step 3: Create Service and Expose it](#)

[Execution](#)

[Step 4: Verify it](#)

Deployment

Deployment Controller changes the actual state to the desired state at a controlled rate. We can create new ReplicaSets, remove existing Deployments and adopt all their resources with new Deployments.

Step 1: Create Deployment File

Create a deployment.yml in a vi editor on your master machine

```
vi deployment.yml
```

Get into the deployment.yml file

```
#the scripts are running from top to bottom and left to right fashion

apiVersion: apps/v1    # in kubernetes all the communication happen through api version
kind: Deployment        # kind of api we are using
metadata:              # using this we provide the name of the api inside cluster
  name: myapp-deployment      # giving name to the deployment
spec:                  # specification
  replicas: 3            # this will create 3 pods inside the deployment
  selector:              # to select resources based on the value of labels of pods
    matchLabels:          # it tells what pods the deployment will apply
      app: myapp3          # label format will be in key value format
  template:              # the template defines reusable content, logic, and parameters
    metadata:
```

```

name: myapp3-pod      # name of the pod
labels:                # maps with the pods in form of key:value pairs
  app: myapp3
spec:                 # indicates that the Pods run one container
  containers:          # stores in form of list
    - name: myapp3-container      # name of the container
      image: piuma/phpsysinfo    # it search/checks for an image with this name
      ports:                   # exposes the Kubernetes service on the specified port within the cluster
        - containerPort: 80       # exposed port 80

```

Step 2: Deploy it & verify it

- When the Deployment file is created . We need to deploy it in our vm with the help of command

```
kubectl apply -f deployment.yml
```

To verify whether the file is deployed or not, we can use can use the following commands

- By getting pod that are available

```
kubectl get pod
```

Output as follows

```

root@master:~# kubectl get pod
              READY   STATUS    RESTARTS   AGE
nginx-deployment-747757dc89-jj5j5   1/1     Running   0          143m

```

- By describing the pod with the hrlp of command

```
kubectl describe pod {name} (nginx-deployment-747757dc89-sr7wx)
```

The describing of pod provides us the following data about that particular pod

```

root@master:~# kubectl describe pod nginx-deployment-747757dc89-sr7wx

Name:           nginx-deployment-747757dc89-sr7wx
Namespace:      default
Priority:       0
Node:           ip-172-31-75-169/172.31.75.169
Start Time:     Wed, 17 Nov 2021 06:46:06 +0000
Labels:         app=nginx
pod-template-hash=747757dc89
Annotations:   <none>
Status:         Running
IP:            10.224.1.15
IPs:
IP:            10.224.1.15
Controlled By: ReplicaSet/nginx-deployment-747757dc89
Containers:
nginx:
Container ID:  containerd://10e2113764d165b8cf25e0b391e47551a13f6e0d519baf312b883150ff616d93
Image:          piuma/phpsysinfo
Image ID:       sha256:8594ee49d529eb313ccb2f26e29e955cf29577ce55834a7cd85586219c486475
Port:          80/TCP
Host Port:     0/TCP
State:         Running
Started:       Wed, 17 Nov 2021 06:46:07 +0000
Ready:          True
Restart Count:  0
Environment:   <none>
Mounts:
/var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-2rkqg (ro)
Conditions:
Type          Status
Initialized   True

```

```

Ready          True
ContainersReady  True
PodScheduled    True
Volumes:
kube-api-access-2rkqg:
Type:           Projected (a volume that contains injected data from multiple sources)
TokenExpirationSeconds: 3607
ConfigMapName:   kube-root-ca.crt
ConfigMapOptional: <nil>
DownwardAPI:    true
QoS Class:     BestEffort
Node-Selectors: <none>
Tolerations:   node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:         <none>

```

the above pod describes that the file has been deployed successfully.

- by the help of replica set

```
kubectl get rs
```

Output as follows

```

root@master:~# kubectl get rs

NAME              DESIRED   CURRENT   READY   AGE
nginx-deployment-66b6c48dd5      0         0         0       3h38m

```

- Describing a particular replicaSet

```
kubectl describe rs {name} (nginx-deployment-66b6c48dd5)
```

Output as follows

```

root@master:~# kubectl describe rs nginx-deployment-66b6c48dd5
Name:           nginx-deployment-66b6c48dd5
Namespace:      default
Selector:       app=nginx,pod-template-hash=66b6c48dd5
Labels:         app=nginx
pod-template-hash=66b6c48dd5
Annotations:   [deployment.kubernetes.io/desired-replicas:] (http://deployment.kubernetes.io/desired-replicas:) 3
[deployment.kubernetes.io/max-replicas:] (http://deployment.kubernetes.io/max-replicas:) 4
[deployment.kubernetes.io/revision:] (http://deployment.kubernetes.io/revision:) 1
Controlled By: Deployment/nginx-deployment
Replicas:      0 current / 0 desired
Pods Status:   0 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
Labels:        app=nginx
pod-template-hash=66b6c48dd5
Containers:
nginx:
Image:         nginx:1.14.2
Port:          80/TCP
Host Port:    0/TCP
Environment:  <none>
Mounts:        <none>
Volumes:      <none>
Events:        <none>

```

Describing of replica set provides the above information including image, port, labels etc...

Service

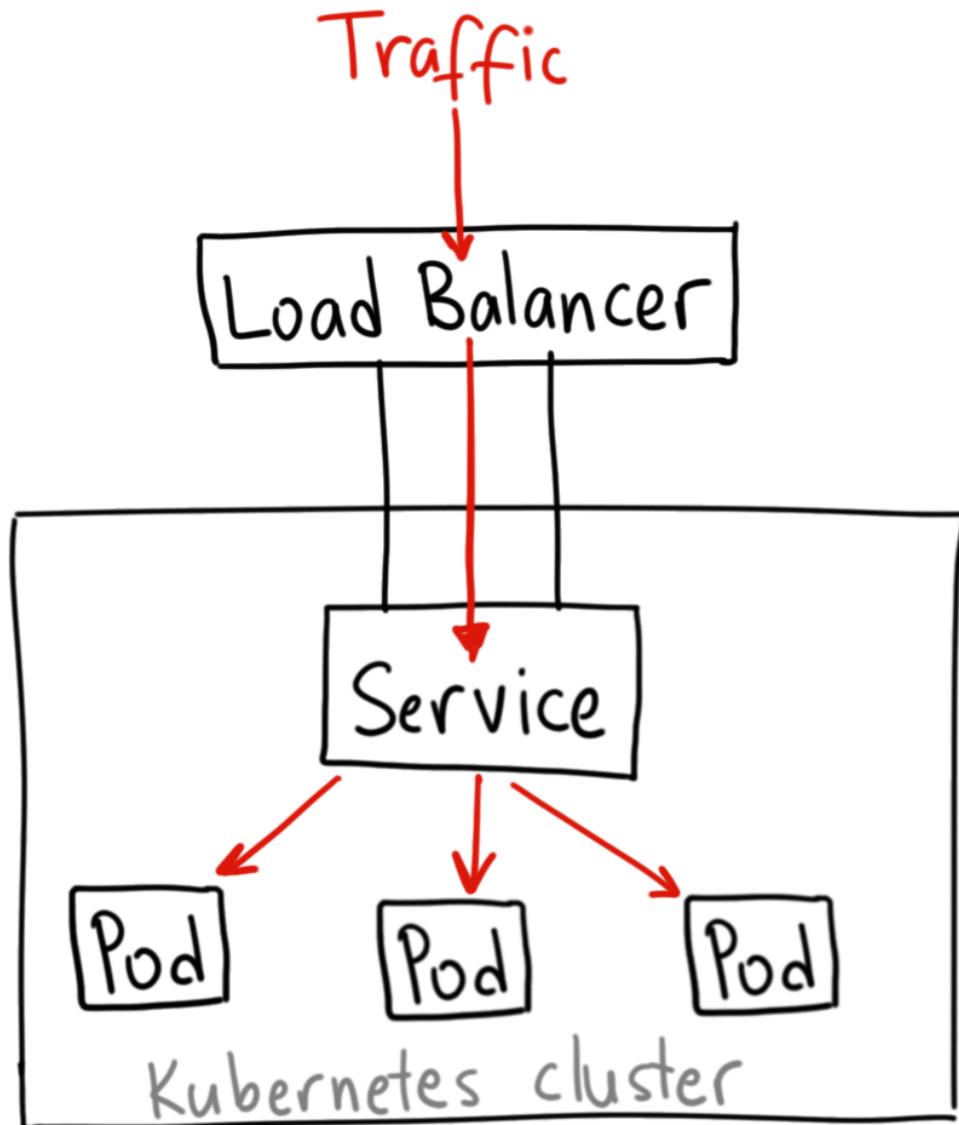
Step 3: Create Service and Expose it

After creating a successful deployment we want to expose our application to the outside world. To do this, we make use of Kubernetes Service.

According to the Kubernetes documentation, Service is an abstract way to expose an application running on a set of pods as a network service.

What does it mean?

When we want to expose our application to the world, we only need to create a new service in Kubernetes. It will automatically map the pods of a deployment to the network and generate its own load balancer.



It does the mapping based on the **pod labels**. After successful execution, we get **endpoints** for each pod which can be seen in the steps for verification.

Execution

In your master machine you can type the following command,

```
kubectl expose deploy myapp-deployment --type=NodePort
```

Let's understand the anatomy.

- `kubectl expose deploy` - this tells the Kubernetes cluster to **expose a deployment**
- `myapp-deployment` - this is the name of our deployment manifest file, this will serve as the specification for our service exposure.
- `--type=NodePort` - this argument tells the service creator to use the worker machine node to expose the deployment.

The expected output after running the command :-

```
service/myapp-deployment exposed
```

This means our application has been successfully exposed.

Step 4: Verify it

After successful execution, we can verify the status of the service with "kubectl get svc"

```
root@ip-172-31-72-71:~/ kubectl get svc
NAME           TYPE      CLUSTER-IP    EXTERNAL-IP   PORT(S)        AGE
kubernetes     ClusterIP  10.96.0.1    <none>        443/TCP       5d2h
my-first-service NodePort  10.101.97.235 <none>        80:31506/TCP  24h
myapp-deployment NodePort  10.102.103.44  <none>        80:30914/TCP  157m
```

Here we can see that our service "myapp-deployment" has been auto-generated from the deployment file. It has been exposed at port 80:30914.

From here we can copy the Port number 30914. This port number should be later attached with the public IP address of the slave/node machine to access the exposed application.

```
root@ip-172-31-72-71:~/capgeminimanifestnov# kubectl get pod -o wide
NAME          READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READIN
my-first-deployment-f5dc549d6-x24fw  1/1    Running   0          23h    10.244.1.12  ip-172-31-78-206  <none>   <none>
my-second-pod   1/1    Running   1 (29h ago)  4d21h  10.244.1.7   ip-172-31-78-206  <none>   <none>
myapp-deployment-7bb646b9dc-2zngx  1/1    Running   0          136m   10.244.1.24  ip-172-31-78-206  <none>   <none>
myapp-deployment-7bb646b9dc-69tm7   1/1    Running   0          136m   10.244.1.23  ip-172-31-78-206  <none>   <none>
myapp-deployment-7bb646b9dc-7pvg3  1/1    Running   0          136m   10.244.1.22  ip-172-31-78-206  <none>   <none>
nginx-deployment-747757dc89-c5tpb  1/1    Running   0          148m   10.244.1.19  ip-172-31-78-206  <none>   <none>
nginx-deployment-747757dc89-ddn8x  1/1    Running   0          148m   10.244.1.20  ip-172-31-78-206  <none>   <none>
nginx-deployment-747757dc89-lh9r4   1/1    Running   0          148m   10.244.1.21  ip-172-31-78-206  <none>   <none>
testpod         1/1    Running   0          28h    10.244.1.10  ip-172-31-78-206  <none>   <none>
```

Here we can see that there are 3 services for myapp-deployment attached with randomly generated string values. They have different IP addresses which are also called **endpoints**.

```
myapp-deployment-7bb646b9dc-2zngx      10.244.1.24  ip-172-31-78-206
myapp-deployment-7bb646b9dc-69tm7      10.244.1.23  ip-172-31-78-206
myapp-deployment-7bb646b9dc-7pvg3      10.244.1.22  ip-172-31-78-206
```

To read complete description of the service, we use the "kubectl describe svc" function.

```
root@ip-172-31-72-71:~/ kubectl describe svc
Name:           myapp-deployment
Namespace:      default
Labels:          <none>
Annotations:    <none>
Selector:       app=myapp3
Type:           NodePort
IP Family Policy: SingleStack
IP Families:   IPv4
IP:              10.102.103.44
IPs:             10.102.103.44
Port:            <unset>  80/TCP
TargetPort:     80/TCP
NodePort:       <unset>  30914/TCP
Endpoints:      10.244.1.22:80,10.244.1.23:80,10.244.1.24:80
Session Affinity: None
External Traffic Policy: Cluster
Events:          <none>
```

We can find various attributes related to our service here. We can see in the Endpoints that all the are added. These are randomly generated. We can also find our NodePort. In case we want to debug the application, we can also read the report in events tab.

Now we have to copy the IP address of the slave/node machine from our AWS console.

The screenshot shows the AWS EC2 Instances page. On the left, there's a sidebar with options like EC2 Dashboard, EC2 Global View, Events, Tags, Limits, Instances (selected), and Images. The main area displays an instance summary for 'i-0630428041100e20f'. It shows the Instance ID, IPv6 address, Private IPv4 DNS, VPC ID, and Subnet ID. The Public IPv4 address '3.227.1.252' is highlighted with a green border and has a tooltip 'Public IPv4 address copied'. The instance state is 'Running' and the type is 't3a.medium'. There's a note about AWS Compute Optimizer finding a user who is not authorized to perform certain actions. The bottom of the screen shows a Windows taskbar with various icons and a system tray indicating the date and time as 17-11-2021 at 14:44.

As shown above copy the IP address and open a new tab in the browser and paste the IP address with the service IP address as shown below

The screenshot shows a Google Chrome browser window. The address bar contains the URL '3.227.1.252:30914'. The main content area displays the Google homepage. Below the address bar, the toolbar includes a search bar with 'Search Google or type a URL' and a microphone icon, followed by several icons for different services like Gmail, Photos, and Drive. The bottom of the screen shows a Windows taskbar with various icons and a system tray indicating the date and time as 17-11-2021 at 14:50.

Now since the web page is opening without any connection error we can say that the services are exposed. The output is as shown below

The screenshot shows a web browser window displaying system information for a Kubernetes deployment named "myapp-deployment-7bb646b9dc-69tm7". The URL is 3.227.1.252:30914/index.php?disp=dynamic. The page has a green header bar and a white content area.

System information : myapp-deployment-7bb646b9dc-69tm7 (10.244.1.23)

Template: phpsysinfo | Language: en

SYSTEM VITAL

Canonical Hostname	myapp-deployment-7bb646b9dc-69tm7
Listening IP	10.244.1.23
Kernel Version	5.11.0-1021-aws (SMP) x86_64
Distro Name	CentOS Linux release 7.1.1503 (Core)
Uptime	1 days 5 hours 18 minutes
Last boot	Tue, 16 Nov 2021 03:56:24 GMT
Current Users	0
Load Averages	0.04 0.11 0.05
System Language	English United States (en_US)
Code Page	UTF-8
Processes	11 (3 running, 8 sleeping)

HARDWARE INFORMATION

Machine	Amazon EC2 t3a.medium, BIOS 1.0 10/16/2017
Processors	AMD EPYC 7571
Processors	AMD EPYC 7571

MEMORY USAGE

Type	Usage	Free	Used	Size
Physical Memory	68%	1.23 GiB	2.56 GiB	3.79 GiB

MOUNTED FILESYSTEMS

Mountpoint	Type	Partition
/	overlay	overlay (nv, relatime, lowerdir=/var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/69/ /var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/68/ upperdir=/var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/96/ /var/lib/containerd/io.containerd.snapshotter.v1.overlayfs/snapshots/88/)
/dev	tmpfs	tmpfs (nv, nosuid, size=65536k, mode=755, inode64)
/dev/shm	tmpfs	shm